



Computation Support in DROP

v6.95 17 August 2025



Java 9 Garbage Collection Algorithms

Overview

1. Role of Java Garbage Collection: *Garbage Collection* (GC) is one of the features behind Java's popularity. Garbage collection is the mechanism used in Java to free unused memory. Essentially, it is tracking down all the objects that are still used and marking the rest as garbage.
2. Automated Memory Management in Java: Java's garbage collection is considered to be an automatic memory management scheme because developers do not have to designate objects as being ready to be deallocated. Garbage collection runs on low-priority threads.
3. Concepts and Algorithms behind GC: This chapter goes through the various concepts behind memory allocation/deallocation, the algorithms that run behind the scene, and the options to customize the behavior.

Object Life Cycle

1. Stages in an Object Life Cycle: A java object's cycle can be seen as composed of being in 3 stages.
2. Object Creation: To create an object, the *new* keyword is generally used, i.e.,

Object obj = new Object();



When an object is created, a specific amount of memory is allocated for storing the object. The amount of memory allocated can differ based on architecture and JVM.

3. Object in use: After creation, the object is used by the application's other objects, i.e., other live objects have references pointing to it. During its usage, objects reside in memory, and may contain references to other objects.
4. Object Destruction: The garbage collection system monitors objects and, if feasible, counts the number of references to each object. When there are no more references to an object, there is no way to get to it with the currently running code, so it makes perfect sense to deallocate the associated memory.

Garbage Collection Algorithms

Object creation is done by the code the developer writes, the frameworks used, and their corresponding features. The developer is not required to deallocate the memory or dereference the objects. This is done automatically at the JVM level by the garbage collector. Since java's inception, there have been many updates to the algorithms which run behind the scene to free the memory. This section examines how they work.

Mark and Sweep

1. Stages of the Algorithm: Mark-and-sweep is a very basic algorithm that runs in two stages:
 - a. Marking live objects - this finds out all the objects that are still alive.



- b. Removing Unreachable Objects - This gets rid of everything else, the supposedly dead and unused objects.
2. Identifying the Garbage Collection Roots: To start with, GC defines some specific objects are *garbage collection roots*, e.g., these include local variables and input parameters of the currently executing methods, active threads, static fields of the loaded classes, and JNI references.
3. Traversing the In-memory Object Graph: The GC traverses the whole object graph in memory, starting from those roots and following references from the roots to other objects. Every object that the GC visits is marked as alive.
4. Stop-the-World Pause: The application threads need to be stopped for the marking to happen as the collector cannot really traverse the graph if it keeps on changing. This is called *stop-the-world pause*.
5. Object Memory Free-up Phase: The second stage is to get rid of unused objects to free up memory. This can be done in a variety of ways, as shown below.
6. Normal Deletion: Normal deletion removes unreferenced objects to free space and leaves referenced objects and pointers. The memory allocator - a kind of hashtable - holds references to blocks of free space where new object can be allocated. This is often referred to as *mark-sweep* allocation.
7. Deletion with Compacting: Only removing unused objects is not efficient, because blocks of free memory are scattered across the storage area and cause *OutOfMemoryError* if the created object is big such that no large enough memory block is available.
8. Motivation behind Compacting Assigned Objects: To solve the above issue, after deleting unreferenced objects, compacting is done on the remaining referenced objects. Here, compacting refers to the process of moving referenced objects together. This makes memory allocation easier, faster, and more reliable. This is often referred to as *mark-sweep-compact* algorithm.



9. Deletion with Copying: This is very similar to the mark and compact approach, as it too relocates all live objects. The important difference is that the target of relocation is a different memory region. This is often referred to as *mark-copy* algorithm.
10. Generational Java Memory Management Scheme: Much of the remaining analysis relies on the structures introduced in Java Memory Management, so it makes sense to review it. In particular, the concepts of young generation, old generation, and permanent generation need to be understood in good detail.

Concurrent Mark-Sweep (CMS) Garbage Collection

1. CMS GC Multi-threaded Garbage Collection: CMS garbage collection is essentially an upgraded mark-and-sweep method. It scans heap memory using multiple threads. It was modified to take advantage of faster systems and has additional performance enhancements.
2. Customized Collection for Young/Old Generation: It attempts to minimize the pauses due to garbage collection by doing most of the garbage collection work *concurrently* with running application threads. It uses the parallel stop-the-world *mark-copy* algorithm in the young generation and the mostly concurrent *mark-sweep* algorithm in the old generation.
3. Invocation of the CMS GC: To invoke CMS GC, the JVM is launched using the argument below: `-XX:+UseConcMarkSweepGC`

CMS GC Optimization Options

Flag	Description
------	-------------



<i>XX: +UseCMSInitiating/OccupancyOnly</i>	Indicates that the invoker wants to solely use occupancy as a criterion for starting a CMS collection operation
<i>XX: CMSInitiating/OccupationFraction = 70</i>	Sets the percentage of CMS generation occupancy to start a CMS generation cycle
<i>XX: CMSTriggerRatio = 70</i>	This is the percentage of <i>MinHeapFreeRatio</i> in the CMS generation that is allocated prior to the start of a CMS cycle
<i>XX: CMSTriggerPermRatio = 90</i>	Sets the <i>MinHeapFreeRatio</i> in the CMS permanent generation that is allocated before starting a CMS collection cycle
<i>XX: CMSWaitDuration = 200</i>	This parameter specifies how long the CMS is allowed to wait for young generation collection
<i>XX: +UseParNewGC</i>	Elect to use the parallel algorithm for young generation collection
<i>XX: +CMSConcurrentMTEnabled</i>	Enables the use of multiple threads for concurrent phases
<i>XX: ConcGCThreads = 2</i>	Sets the number of parallel threads for concurrent phases
<i>XX: +CMSIncrementalMode</i>	Enables the incremental CMS (iCMS) mode
<i>XX: +CMSClassUnloadingEnabled</i>	If this is not enabled, CMS will not clean permanent space
<i>XX: +ExplicitGCInvokes/Concurrent</i>	This allows <i>System.gc()</i> to trigger concurrent collection instead of a full garbage collection cycle

Serial Garbage Collection

1. Mechanism of Serial Garbage Collection: This algorithm uses *mark-copy* for young generation and *mark-sweep-compact* for old generation. It works on a single thread. When executing, it freezes all other threads until the garbage collection operations have concluded.



2. Limitations of Serial Garbage Collection: Due to the thread freezing nature of serial garbage collection, it is only feasible for very small programs.
3. Invocation of the Serial Garbage Collector: The JVM arguments below are used to invoke the serial garbage collector: `-XX:+UseSerialGC`

Parallel Garbage Collection

1. Mechanism of Parallel Garbage Collection: Similar to serial GC, parallel garbage collection uses *mark-copy* in the young generation and *mark-sweep-compact* in the old generation. Multiple concurrent threads are used for the marking and the copying/compacting phases. The number of threads may also be configured using the `-XX:ParallelGCThreads = N` option.
2. Improvements achieved by Parallel GC: Parallel garbage collector is suitable on multi-core machines in cases where the primary goal is to increase throughput by efficient usage of existing system resources. Using this approach, GC cycle times can be considerably reduced.
3. Default Collector - Parallel vs G1: Till Java 8, parallel GC was used as the default garbage collector. Java 9 onwards uses G1 as the default garbage collector on 32-bit and 64-bit server configurations.
4. Invoking the Parallel GC Scheme: The JVM arguments below are used to invoke the parallel garbage collector: `-XX:+UseParallelGC`

G1 Garbage Collection

1. G1 - Enhanced CMS Collector: The G1 - or Garbage First - garbage collector was available in Java 7 and is designated to be the long-term replacement for the CMS collector. The G1 collector is a parallel, concurrent, and incrementally compacting low-pause garbage collector.



2. Parallel Collection over Smaller Sub-regions: G1 garbage collection approach involves segmenting the memory heap into multiple small regions, typically 2048. Each region is marked as either young generation - further divided into Eden regions and Survivor regions - or old generation. This allows GC to avoid collecting the entire heap at one, and instead approach the problem incrementally. It means that only a subset of these regions is considered at a time.
3. Maintaining Collection Statistics over Sub-regions: G1 keeps track of the amount of live data each region contains. The information is used in determining the regions that contain the most garbage, so they are collected first. This is why the scheme is named *garbage-first collection*.
4. Tweaking Stop-the-World Durations: Just like other algorithms, unfortunately, the compacting operation takes place using the *stop-the-world* approach. However, specific performance goals can be set as per the design objectives.
5. No Guarantee in Duration Limitations: The pause duration can be configured, e.g., no more than 10 milliseconds in any given second. Garbage first GC will do its best to meet this goal with high probability, but not with certainty, since hard real-time is not guaranteed due to OS level thread management.
5. Invoking the G1 Garbage Collector: In Java 7/8 machines, the JVM argument to invoke the G1 garbage collector is `-XX:+UseG1GC`

G1 Optimization Options

Flag	Description
------	-------------



<i>XX:G1HeapRegionSize = 16m</i>	Size of the Heap region. The value will be a power of 2 and can range from 1 MB to 32 MB. The goal is to have around 2048 regions based on the minimum Java heap size.
<i>XX:MaxGCPauseMillis = 200</i>	Sets the target value for the desired maximum pause time. The default value is 200 milliseconds. The specified value does not adapt to the heap size.
<i>XX:G1ReservePercent = 5</i>	This determines the minimum reserve in the heap.
<i>XX:G1ConfidencePercent = 5</i>	This is the confidence coefficient pause prediction heuristics.
<i>XX:GCPauseIntervalMillis = 5</i>	This is the pause interval time slice per MMU in milliseconds.

G1 Configuration Flags

Flag	Description
<i>-Xms2048m -Xmx3g</i>	Sets the initial and the maximum heap size - young space plus tenured space.
<i>XX:+DisableExplicitGC</i>	This will cause the JVM to ignore any <i>System.gc()</i> method invocation by an application.
<i>XX:+UseGCOverheadLimit</i>	This policy is used to limit the time spent in garbage collection before an <i>OutOfMemory</i> error is thrown.
<i>XX:GCTimeLimit = 95</i>	This limits the proportion of time spent in garbage collection before an <i>OutOfMemory</i> error is thrown.
<i>XX:GCHeapFreeLimit = 5</i>	This sets the minimum percentage of free space after a full garbage collection before an <i>OutOfMemory</i> error is thrown. This is used with <i>GCTimeLimit</i> .
<i>XX:InitialHeapSize = 3g</i>	Sets the initial heap size - young space plus tenured space.
<i>XX:MaxHeapSize = 3g</i>	Sets the maximum heap size - young space plus tenured space.
<i>XX:NewSize = 128m</i>	Sets the initial size of young space.
<i>XX:MaxNewSize = 128m</i>	Sets the maximum size of young space.



<i>XX:SurvivorRatio = 15</i>	Sets the size of the single survivor space as a portion of the Eden space size.
<i>XX:PermSize = 512m</i>	Sets the initial size of the permanent size.
<i>XX:MaxPermSize = 512m</i>	Sets the maximum size of the permanent size.
<i>-Xss512k</i>	Sets the size of the area dedicated to each thread in bytes.

G1 Logging Flags

Flag	Description
<i>-verbose:gc</i> OR <i>-XX:+PrintGC</i>	This prints the basic garbage collection information.
<i>-XX:+PrintGCDetails</i>	This will print more detailed garbage collection information.
<i>XX:+PrintGCTimestamps</i>	This prints timestamps for each garbage collection event. The seconds are sequential and begin from the JVM start time.
<i>XX:+PrintGCDateStamps</i>	This prints the date-stamps for each garbage collection event.
<i>-Xloggc:</i>	Using this directs the garbage collection output to a file instead of the console.
<i>XX:+Print\TenuringDistribution</i>	This prints detailed information regarding young space following each collection cycle.
<i>-XX:+PrintTLAB</i>	This flag prints TLAB collection statistics.
<i>-XX:+PrintReferenceGC</i>	This flag prints the time for reference processing (that is, weak, soft, and so on) during stop-the-world pauses.
<i>-XX:+HeapDump\OnOutOfMemoryError</i>	This flag creates heap dump file in an out-of-memory condition.



Processes and Threads

Overview

1. Java Concurrent Programming using Threads: In concurrent programming, there are two basic units of execution: *processes* and *threads*. In the Java programming language, concurrent programming is mostly concerned with threads. However, processes are also important.
2. Time Slicing Based on Concurrent Threads: A computer system normally has many active processes and threads. This is true even in systems that only have a single execution core, and thus have only one thread executing at any given moment. Processing time for a single core is shared among processes and threads through an OS feature called *time slicing*.
3. Concurrency on Single/Multiple Cores: It is becoming more and more common for computer systems to have multiple processors or processors with multiple execution cores. This greatly enhances a system's capacity for concurrent execution of processes and threads, but concurrency is possible even on simple systems, without multiple processors or execution cores.

Processes

1. Self-contained Unit of Execution: Each process has a self-contained execution environment. A process has a complete, private set of basic run-time resources; in particular, each process has its own memory space.
2. Mechanisms of Inter-Process-Communication: Processes are often seen as synonymous with programs or applications. However, what the user sees as



a single application may actually be in fact a set of cooperating processes. To facilitate communications between processes, most operating systems support *inter-process communication* (IPC) resources, such as pipes and sockets. IPC is used not just for communication between processes in the same system, but processes on different systems.

3. Process Creation Scheme in Java: Most implementations of the Java VM run as a single process. A Java application can create additional processes using a *ProcessBuilder* object.

Threads

1. Execution Environment Provided by Threads: Threads are sometimes called *lightweight* processes. Both processes and threads provide an execution environment, but creating a new thread requires fewer resources than creating a new process.
2. Resources Shared by Processes' Threads: Threads exist within a process - every process has at least one. Threads share the processes' resources, including memory and open files. This makes for efficient, but potentially problematic, communication.
3. Process' Main and Supporting Threads: Multi-threaded execution is an essential feature of the Java platform. Every application has at least one thread - or several, if one counts *system* threads that do things like memory management and signal handling. But, from the application programmer's point of view, one starts with just a single thread, called the *main thread*. This thread has the ability to create additional threads, as is demonstrated later.



Thread Objects

Overview

Each thread is associated with an instance of the class *Thread*. There are two basic strategies for using *Thread* objects to create a concurrent application:

- a. To directly control thread creation and management, simply instantiate *Thread* each time the application needs to initiate an asynchronous task.
- b. To abstract thread management from the rest of the application, the application's tasks are passed to an *executor*.

This chapter documents the use of *Thread* objects. Executors are discussed along with other high-level concurrency objects.

Defining and Starting a Thread

1. Provide a *Runnable* Object: An application that creates an instance of *Thread* must provide the code that will run in that method. There are two ways to do this: First, the *Runnable* interface defines a single method, *run*, meant to contain the code executed in the thread. The *Runnable* object is passed to the *Thread* constructor.
2. Subclass *Thread*: The *Thread* class itself implements *Runnable*, though its *run* method does nothing. An application can subclass *Thread*, providing its own implementation of *run*. Both the idioms above will need to invoke *Thread.invoke* in order to start a new thread.
3. Advantages of Implementing *Runnable* Interface: The first idiom, which employs a *Runnable* object, is more general, because the *Runnable* object



can subclass a class other than *Thread*. The second idiom is easier to use in simple applications, but is limited by the fact that the task class must be a descendant of *Thread*. This chapter focusses on the first approach, which separates the *Runnable* task from the *Thread* object that executes the task. Not only is this approach more flexible, it is also applicable to the high-level thread management APIs covered later.

4. Functionality Provided by *Thread* Class: The *Thread* class defines a number of methods useful for thread management. These include *static* methods, which provide information, or affect the status of, the thread invoking the method. The other methods are invoked from other threads involved in managing the thread and the *Thread* object. The next sections examine some of these methods.

Pausing Execution with *sleep*

1. Setting up the *sleep* Time: *Thread.sleep* causes the current thread to suspend execution for a specified period. This is an efficient means of making the processor threads available to the threads of an application or other applications that may be running on a computer system. The *sleep* method can be used for pacing, and waiting for another thread with duties that are understood to have time requirements.
2. Time Guarantees Provided by *sleep*: Two overloaded versions of *sleep* are provided: one specifies the sleep time to the millisecond and one specifies the sleep time to the nanosecond. However, these sleep times are not guaranteed to be precise, because they are limited by the facilities provided by the underlying OS. Also, the sleep period can be terminated by interrupts as seen in a later section. In any case, one cannot assume that invoking *sleep* will suspend the thread for precisely the time period specified.



3. Handling *InterruptedException* Thrown in *sleep*: The method calling *sleep* declares that it *throws InterruptedException*. This is an exception that *sleep* throws when another thread interrupts the current thread while *sleep* is active. If an application has not defined another thread to cause an interrupt, it need not bother catching *InterruptedException*.

Interrupts

1. Purpose behind the Interrupt Construct: An *interrupt* is an indication to a thread that it should stop what it is doing and do something else. It is up to the developer to decide exactly how a thread responds to an interrupt, but it is common for the thread to terminate. This is the usage emphasized in this section.
2. Mechanism for Sending an Interrupt: A thread sends an interrupt by invoking *interrupt* on the *Thread* object of the thread to be interrupted. For the interrupt mechanism to work correctly, the interrupted thread must support its own interruption.

Supporting Interruption

1. Thread Supporting its own Interruption: How does a thread support its own interruption? This depends on what it is currently doing. If the thread is frequently invoking methods that throw *InterruptedException*, it simply returns from the *run* method after it catches the exception. Many methods that throw *InterruptedException*, such as *sleep*, are designed to cancel their current operation and return immediately when an interrupt is received.



2. Methods that don't throw *InterruptedException*: What if a thread goes a long time without invoking a method that throw *InterruptedException*? Then it must periodically invoke *Thread.interrupted*, which return *true* if an interrupt has been received.
3. Explicitly Checking for Triggered Interrupts: In simple applications, the code can simply test for a triggered interrupt and exit the thread if one has been received. In more complex applications, it might make more sense to throw an *InterruptedException*. This would allow interrupt handling code to be centralized in a *catch* clause.

The Interrupt Status Flag

1. Usage Caveats Underlying Interrupt Mechanism: The interrupt mechanism is implemented using an internal flag known as *interrupt status*. Invoking *Thread.interrupt* sets this flag. When a thread checks for an interrupt by invoking the static method *Thread.interrupted*, the interrupt status is cleared. The non-static *isInterrupted* method, which is used by one thread to query the interrupt status of another, does not change the interrupt status flag.
2. Clearing Interrupt Status after throwing *InterruptedException*: By convention, any method that exits by throwing an *InterruptedException* clears the interrupt status when it does so. However, it is always possible that the interrupt status will be immediately set again, by another thread invoking *interrupt*.

Joins



The *join* method allows one thread to wait for the completion of another. If *t* is a *Thread* object whose thread is currently executing, *t.join()* causes the current thread to pause execution until *t*'s thread terminates. Overloads of *join* allow a developer to specify a waiting period. However, as with *sleep*, *join* is dependent on the OS for timing, so one cannot assume that *join* will wait exactly as long as specified. Further, like *sleep*, *join* responds to an interrupt by exiting with an *InterruptedException*.



Google Guice

Overview

1. Google Guice Dependency Injection Suite: *Google Guice* (Wikipedia (2020))
is an open-source software framework for the Java platform released by Google under the Apache license. It provides support for dependency injection using annotations to configure Java objects (Yuan (2008)).
Dependency injection is a design pattern whose core principle is to separate behavior from dependency resolution.
2. @Inject Annotation for Dependency Injection: Guice allows implementation classes to be programmatically bound to an interface, then injected into constructors, methods, or fields using an @Inject annotation. When more than one implementation of the same interface is needed, the user can create custom annotations that identify an implementation, then use that annotation when injecting it.

Motivation

1. Assembling Components of an Application: Wiring everything together is a tedious part of application development. There are several approaches to connect data, services, and presentation classes to one another. To contrast these approaches, the rest of this chapter sketches out the code for a pizza ordering website.
2. The Pizza Ordering Billing Service:



public interface BillingService

```
{  
    /*  
    * Attempts to charge an order to the credit card. Both successful and  
    * failed transactions can be recorded.  
    *  
    * @return A receipt of the transaction. If the charge was successful,  
    *         the receipt will be successful. Otherwise, the receipt will  
    *         contain a note describing why the charge failed.  
    */  
  
    Receipt chargeOrder (  
        final PizzaOrder order,  
        final CreditCard creditCard);  
}
```

3. Unit Tests for the Service: Along with the implementation, unit tests will also be written.

Direct Constructor Calls

1. Credit Card Processor/Transaction Logger Constructor: Here is how the code looks like on the construction of a new credit card processor and a new transaction logger.
2. Hard Construction Billing Service Implementation:

public class RealBillingService implements BillingService

```
{  
    Receipt chargeOrder (  

```



```
final PizzaOrder order,
final CreditCard creditCard)
{
    CreditCardProcessor processor =
        new PaypalCreditCardProcessor();
    TransactionLog transactionLog =
        new DatabaseTransactionLog();

    try
    {
        ChargeResult result = processor.charge (
            creditCard,
            order.amount()
        );
        transactionLog.logChargeResult (
            result
        );
        return result.successful() ?
            Receipt.successfulCharge (
                order.amount()
            ) ? Receipt.declinedCharge (
                result.declineMessage()
            )
        );
    }
    catch (Exception e)
    {
        transactionLog.logConnectException (
            e
        );
    }
}
```



```
        return Receipt.successfulCharge (  
            e.getMessage()  
        );  
    }  
}  
}
```

4. Problems with Hard Constructor Schemes: The above code poses problems for modularity and testability. The direct, compile time dependency on the real credit card processor means that testing the code will charge a real credit card! It is also awkward to test what happens when the charge is declined or when the service is unavailable.

Factories

1. De-coupling Usage from Implementation: A factory class decouples the client and the implementing class. A simple factory class uses static methods to get and set mock interfaces. A factory is implemented with some boiler-plate code.
2. Hard Construction Billing Service Implementation:

```
public class CreditCardProcessorFactory  
{  
    private static CreditCardProcessor _processor = null;  
    public static void SetInstance (  
        final CreditCardProcessor processor)  
    {  
        _processor = processor;  
    }  
}
```



```
public static CreditCardProcessor GetInstance()
{
    if (null == _processor)
    {
        _processor = new SquareCreditCardProcessor();
    }
    return _processor;
}
}
```

3. Replacing Constructor with Factory Lookups:

```
public class RealBillingService implements BillingService
{
    Receipt chargeOrder (
        final PizzaOrder order,
        final CreditCard creditCard)
    {
        CreditCardProcessor processor =
            CreditCardProcessorFactory.GetInstance();
        TransactionLog transactionLog =
            TransactionLogFactory.GetInstance();

        try
        {
            ChargeResult result = processor.charge (
                creditCard,
                order.amount()
            );
            transactionLog.logChargeResult (
```



```
        result
    );
    return result.successful() ?
        Receipt.successfulCharge (
            order.amount()
        ) ? Receipt.declinedCharge (
            result.declineMessage()
        )
    );
}
catch (Exception e)
{
    transactionLog.logConnectException (
        e
    );
    return Receipt.successfulCharge (
        e.getMessage()
    );
}
}
```

3. Advantages in Building Unit Tests: This factory makes it possible to write proper unit tests:

```
public class RealBillingServiceTest extends TestCase
{
    private final PizzaOrder _order = new PizzaOrder (
        100
    );
}
```



```
private final CreditCard _creditCard = new CreditCard (
    "1234",
    11,
    2010
);
private final InMemoryTransactionLog _transactionLog =
    new InMemoryTransactionLog ();
private final FakeCreditCardProcessor _processor =
    new FakeCreditCardProcessor();
@Override public void setUp()
{
    TransactionLogFactory.SetInstance (
        _transactionLog
    );
    CreditCardProcessorFactory.SetInstance (
        _processor
    );
}
@Override public void tearDown()
{
    TransactionLogFactory.SetInstance (
        null
    );
    CreditCardProcessorFactory.SetInstance (
        null
    );
}
public void testSuccessfulCharge()
{
    RealBillingService billingService = new RealBillingService();
```




```
Receipt receipt = billingService.chargeOrder (
    order,
    creditCard
);
assertTrue (
    receipt.hasSuccessfulCharge()
);
assertEquals (
    100,
    receipt.amountOfCharge()
);
assertEquals (
    creditCard,
    _processor.cardOfOnlyCharge()
);
assertEquals (
    100,
    _processor.amountOfOnlyCharge()
);
assertTrue (
    _transactionLog.wasSuccessLogged()
);
}
}
```

4. Factor Problem #1 - Global Instances: The above code is clumsy. A global variable holds the mock implementation, so one needs to be careful about setting it up and tearing it down. Should *tearDown* fail, the global variable continues pointing to the test instance. This could cause problems for other tests. It also prevents from running multiple tests in parallel.



5. Factor Problem #2 - Embedded Dependencies: But the biggest problem is that the dependencies are *hidden in the code*. If a dependency is added to *CreditCardFraudTracker*, the tests will have to be re-run to find out which ones will break. If one forgets to initialize a factory for a production service, one doesn't find out until a change is attempted. As the application grows, babysitting factories becomes a growing drain on productivity.

Dependency Injection

1. Principle behind Dependency Injection: Like factories, dependency injection is just a dependency pattern. The core principle is to *separate behavior from dependency resolution*. In the example below, *RealBillingService* is not responsible for looking up the *TransactionLog* and *CreditCardProcessor*. Instead they are passed in as constructor parameters.
2. Parameterized Construction Billing Service Implementation:

```
public class RealBillingService implements BillingService
{
    private final TransactionLog _transactionLog;
    private final CreditCardProcessor _processor;
    public RealBillingService (
        final CreditCardProcessor processor,
        final TransactionLog transactionLog)
    {
        _processor = processor;
        _transactionLog = transactionLog;
    }
    Receipt chargeOrder (
        final PizzaOrder order,
```



```
        final CreditCard creditCard)
    {
        try
        {
            ChargeResult result = _processor.charge (
                creditCard,
                order.amount()
            );
            _transactionLog.logChargeResult (
                result
            );
            return result.successful() ?
                Receipt.successfulCharge (
                    order.amount()
                ) ? Receipt.declinedCharge (
                    result.declineMessage()
                )
            );
        }
        catch (Exception e)
        {
            _transactionLog.logConnectException (
                e
            );
            return Receipt.successfulCharge (
                e.getMessage()
            );
        }
    }
}
```



3. Eliminating Set Up/Tear Down: There is no need for factories, and the test cases can be simplified by removing the *setUp* and *tearDown* boiler-plate.
4. Simplified Billing Service Unit Test:

```
public class RealBillingServiceTest extends TestCase
{
    private final PizzaOrder _order = new PizzaOrder (
        100
    );
    private final CreditCard _creditCard = new CreditCard (
        "1234",
        11,
        2010
    );
    private final InMemoryTransactionLog _transactionLog =
        new InMemoryTransactionLog ();
    private final FakeCreditCardProcessor _processor =
        new FakeCreditCardProcessor();
    public void testSuccessfulCharge()
    {
        RealBillingService billingService = new RealBillingService (
            _processor,
            _transactionLog
        );
        Receipt receipt = billingService.chargeOrder (
            order,
            creditCard
        );
        assertTrue (
```



```
        receipt.hasSuccessfulCharge()
    );
    assertEquals (
        100,
        receipt.amountOfCharge()
    );
    assertEquals (
        creditCard,
        _processor.cardOfOnlyCharge()
    );
    assertEquals (
        100,
        _processor.amountOfOnlyCharge()
    );
    assertTrue (
        _transactionLog.wasSuccessLogged()
    );
}
}
```

5. Compile Time Dependent Class Detection: Now, whenever dependencies are added or removed, the compiler reminds what tests need to be fixed. This dependency is *exposed in the API signature*.
6. Recursive Application of DI Pattern: Unfortunately, the clients of *BillingService* now need to look up its dependencies. Some of these can be fixed by applying the pattern again. Classes the depend on it can accept a *BillingService* in their constructor. For top-level classes, it is useful to have a framework. Otherwise, the dependencies need to be recursively constructed when one needs to use a service:



```
public static void main (  
    final String[] argumentArray)  
{  
    CreditCardProcessor processor = new PaypalCreditCardProcessor();  
    TransactionLog transactionLog = new DatabaseTransactionLog();  
    BillingService billingService = new RealBillingService (  
        processor,  
        transactionLog  
    );  
    ...  
}
```

Dependency Injection with Guice

1. Using Guice for Dependency Injection: The dependency injection pattern leads to code that is modular and testable, and Guice makes it easy to write. To use Guice in the billing service, one needs to tell it how to map the interfaces to implementations. This configuration is done in the Guice module, which is any Java class that implements the *Module* interface.
2. Guice - Binding Implementations to Interfaces:

```
public class BillingModule implements AbstractModule  
{  
    @Override protected void configure()  
    {  
        bind (  
            TransactionLog.class  
        ).to (  
            DatabaseTransactionLog.class  
        )  
    }  
}
```



```
);  
bind (  
    CreditCardProcessor.class  
).to (  
    PaypalCreditCardProcessor.class  
);  
bind (  
    BillingService.class  
).to (  
    RealBillingService.class  
);  
}  
}
```

3. Compile Time Dependent Class Detection: Adding *@Inject* to *RealBillingService*'s constructor directs Guice to use it. Guice will inspect the annotated constructor, and lookup values for each parameter.
4. Parameterized Construction Billing Service Implementation:

```
public class RealBillingService implements BillingService  
{  
    private final TransactionLog _transactionLog;  
    private final CreditCardProcessor _processor;  
    @Inject public RealBillingService (  
        final CreditCardProcessor processor,  
        final TransactionLog transactionLog)  
    {  
        _processor = processor;  
        _transactionLog = transactionLog;  
    }  
}
```



```
Receipt chargeOrder (  
    final PizzaOrder order,  
    final CreditCard creditCard)  
{  
    try  
    {  
        ChargeResult result = _processor.charge (  
            creditCard,  
            order.amount()  
        );  
        _transactionLog.logChargeResult (  
            result  
        );  
        return result.successful() ?  
            Receipt.successfulCharge (  
                order.amount()  
            ) ? Receipt.declinedCharge (  
                result.declineMessage()  
            )  
            ;  
    }  
    catch (Exception e)  
    {  
        _transactionLog.logConnectException (  
            e  
        );  
        return Receipt.successfulCharge (  
            e.getMessage()  
        );  
    }  
}
```




```
}  
}
```

5. Injecting an Arbitrary Bound Instance: Finally, putting it all together, the *Injector* can be used to get an instance of any of the bound classes.

```
public static void main (  
    final String[] argumentArray)  
{  
    Injector injector = Guice.createInjector (  
        new BillingModule()  
    );  
    BillingService billingService = injector.getInstance (  
        BillingService.class  
    );  
}
```

References

- Yuan, M. (2008): [Guice \(Google\)](#)
- Wikipedia (2020): [Google Guice](#)



Terraform

Abstract

Terraform enables safe and predictable creation, change, and improvement in the infrastructure. It is an open-source tool that codifies APIs into declarative configuration files that can be shared among team members and treated as code, edited, reviewed, and versioned.

Overview

1. Design Objective Behind Terraform Configurator: *Terraform* is an open-source infrastructure-as-code software tool created by HashiCorp. Users define and provision data center infrastructure using a declarative configuration language known as HashiCorp Configuration Language - HCL - or optionally JSON (Wikipedia (2020)).
2. Declarative Management of External Resources: Terraform manages external resource such as public cloud infrastructure, private cloud infrastructure, network appliances, software-as-a-service, and platform-as-a-service with *providers*. HashiCorp maintains an extensive list of official providers, and can also integrate with community-developed providers.
3. Interaction with the Terraform Providers: Users can interact with the Terraform providers by declaring resources or by calling data sources. Rather than using imperative commands to provision resources, Terraform uses declarative configuration to specify the desired final state. Once the



user invokes Terraform on a given resource, Terraform will perform CRUD actions on the user's behalf to achieve the desired state.

4. Modularization of Infrastructure Management: Infrastructure-as-code can be written as modules, promoting reusability and maintainability.
5. List of Infrastructure Providers Supported: Terraform supports a number of cloud infrastructure providers such as Amazon Web Services, Microsoft Azure, IBM Cloud, Google Cloud Platform, Digital Ocean, Oracle Cloud Infrastructure, VMWare vSphere, and OpenStack (Somwanshi (2015), Turnbull (2016), Brikman (2017)).
6. Enhancements to the Terraform Suite: Since 2017, HashiCorp has also started supporting a Terraform Module Registry. In 2019, Terraform introduced the paid version called Terraform Enterprise for larger organizations.

Components

1. Role of the Terraform Tool: Terraform can manage existing and popular service providers as well as custom in-house solutions.
2. Infrastructure as Code: Infrastructure is described using a high-level configuration syntax. This allows a blueprint of the data center to be versioned and treated as any other code. Additionally, infrastructure can be shared and reused.
3. Main Commands of the Terraform Suite: Terraform has four major commands: terraform init, terraform plan, terraform apply, and terraform destroy.
4. Execution Plans: Terraform has a *planning* step where it generates an *execution plan*. The execution plan shows what Terraform will do when *apply* is called. This avoids surprises when Terraform manipulates infrastructure.



5. Resource Graph: Terraform builds a graph of all resources, and parallelizes the creation and the modification of any non-dependent resources. Because of this, Terraform builds infrastructure as efficiently as possible, and operators get insight into the dependencies in their infrastructure.
6. Change Automation: Complex changesets can be applied to the infrastructure with minimal human interaction. With the previously mentioned execution plan and resource graph, one knows exactly what Terraform will change and in what order, avoiding many possible human errors.

References

- Brikman, Y. (2017): *Terraform: Writing Infrastructure as Code* O'Reilly
- Somwanshi, S. (2015): [Choosing the Right Tool to Provision AWS Infrastructure](#)
- Turnbull, J. (2016): *The Terraform Book, 151st Edition* Turnbull Press
- Wikipedia (2020): [Terraform \(software\)](#)



Terraform Use Cases

Overview

This chapter lists some concrete use cases for Terraform, but the possible use cases are much broader than what is covered here. Due to its extensible nature, providers and provisioners can be added to further extend Terraform's ability to manipulate resources.

Heroku App Setup

1. PaaS for Hosting Web Apps: Heroku is a popular PaaS for hosting web apps. Developers can create an app, and then attach add-ons, such as a database, or e-mail provider. One of the best features is the ability to elastically scale the number of dynos or workers. However, most non-trivial applications quickly need many add-ons and external services.
2. Codification of a Heroku Setup: Terraform can be used to codify the setup needed for a Heroku application, ensuring that all required add-ons are available, but it can go even further: configuring DNSimple to set a CNAME, or setting up Cloudflare as a CDN for the app. Best of all, Terraform can do all this in under 30 seconds without using a web interface.

Multi-Tier Applications



1. Motivation behind Multi-tier Applications: A very common pattern is the N-tier architecture. The most common 2-tier architecture is a pool of web servers that use a database tier. Additional tiers get added for API servers, caching servers, routing meshes, etc. This pattern is used because the tiers can be scaled independently and provide a separation of concerns.
2. Setting Terraform to Handle Tiers: Terraform is an ideal tool for building and managing these infrastructures. Each tier can be described as a collection of resources, and the dependencies between each tier are handled automatically. Terraform will ensure that the database tier is available before the web servers are started and that the load balancers are aware of the web modules.
3. Scaling Tiers Independently using Terraform: Each tier can then be scaled using Terraform by modifying single *count* configuration value. Because creation and provisioning of a resource is codified and automated, elastically scaling with load becomes trivial.

Self-Service Clusters

1. Self-Service Infrastructure Cluster Blueprints: At a certain organizational size, it becomes very challenging for a centralized operations team to manage a large and growing infrastructure. Instead it becomes more attractive to make *self-service* infrastructure, allowing product teams to manage their own infrastructure using tooling provided by the central operations team.
2. Transmission/Usage of Blueprint Templates: Using Terraform, the knowledge of how to build and scale a service can be codified in a configuration. Terraform configurations can be shared within an organization enabling customer teams to use the configuration as a black-box and use Terraform as a tool to manage their services.



Software Demos

1. Realistic Virtualized Environment for Demos: Modern software is increasingly networked and distributed. Although tools like Vagrant (<https://www.vagrantup.com/>) exist to build virtualized environment for demos, it is still very challenging to demo software on real infrastructure which more closely matches production environments.
2. Terraform Configuration to the Rescue: Software developers can provide a Terraform configuration to create, provision, and bootstrap a demo on cloud providers like AWS. This allows end users to easily demo software on their own infrastructure, and even enables tweaking parameters like the cluster size to more rigorously test tools at any stage.

Disposable Environments

1. Partitioning into Production/Staging/QA: It is a common practice to have both a production and staging/QA environment. These environments are smaller clones of their production counterparts, but are used to test new applications before releasing them into production. As the production environment grows larger and more complex, it becomes increasingly onerous to maintain an up-to-date staging environment.
2. Codification of Production Environment Setup: Using Terraform, the production environment can be codified and then shared with staging, QA, or dev. These configurations can be used to rapidly spin up new environments, to test in, and then be easily disposed of. Terraform can help tame the difficulty of maintaining parallel environments, and makes it practical to elastically crate and destroy them.



Software Defined Networking

1. Advantages of Software Defined Network: Software Defined Networking (SDN) is becoming increasingly prevalent in the datacenter, as it provides more control to operators and developers and allows the network to better support the applications running on top. Most SDN implementations have a control layer and an infrastructure layer.
2. Codifying a Software Defined Network: Terraform can be used to codify the configuration for Software Defined Networks. This configuration can then be used by Terraform to automatically setup and modify settings by interfacing a control layer. This allows configuration to be versioned and changes to be automated. As an example, AWS VPC - <https://aws.amazon.com> - is one of the most commonly used SDN implementations, and can be configured by Terraform - `/docs/providers/aws/r/vpc.html`

Resource Schedulers

1. Large-Scale Dynamic Resource Scheduling: In large-scale infrastructures, static assignment of applications to machines become increasingly challenging. To solve that problem, there are a number of schedulers like Borg, Mesos, YARN, and Kubernetes. These can be used to dynamically schedule Docker containers, Hadoop, Spark, and many other software tools.
2. Terraform Configuration of Resource Schedulers: Terraform is not limited to physical providers like AWS. Resource schedulers are treated like providers, enabling Terraform to request resources from them. This allows Terraform to be used in layers; to setup the physical infrastructure running the layers as well as provisioning onto the scheduled grid.



Multi-Cloud Deployment

1. Advantages of Multi-Cloud Deployment: It is often attractive to spread infrastructure across multiple clouds to increase fault-tolerance. By using only a single region or cloud provider, fault tolerance is limited by the availability of that provider. Having a multi-cloud deployment allows for a more graceful recovery from the loss of a region or an entire provider.
2. Terraform Based Multi-Cloud Deployment: Realizing multi-cloud deployments can be very challenging as many existing tools for infrastructure management are cloud-specific. Terraform is cloud agnostic and allows a single configuration to be used to manage multiple providers, and even to handle cross-cloud dependencies. This simplifies management and orchestration, helping operators build large-scale multi-cloud infrastructure.



Terraform vs. Other Software

Overview

1. Abstraction across Resources and Providers: Terraform provides flexible abstraction of resources and providers. This model allows for representing everything from physical hardware, virtual machines, and containers, to email and DNS servers.
2. Overlapping with the Universe of Existing Tools: Because of this flexibility, Terraform can be used to solve many different problems. This means that there are a number of existing tools that overlap with the capabilities of Terraform.
3. Terraform Coexistence in this Universe: This chapter compares Terraform to a number of these tools, but it should be noted that Terraform is not mutually exclusive with other systems. It can be used to manage a single application, or an entire datacenter.

Terraform vs. Chef, Puppet, etc.

1. Purpose of Configuration Management Tools: Configuration management tools install and manage software on a machine that already exists. Terraform is not a configuration management tool, and it allows existing tooling to focus on their strengths: bootstrapping and initializing resources.
2. Coexistence in a Terraform Universe: Using provisioners, Terraform enables any configuration management to be used to setup a resource once it has been created. Terraform focuses on the higher-level abstraction of the



datacenter and associated services, without sacrificing the ability of the configuration management tools to do what they do best. It also embraces the same codification that is responsible for the success of these tools, making entire infrastructure deployments easy and reliable.

Terraform vs. Cloud Formation, Heat etc.

1. File-based Infrastructure Configuration Specification: Tools like Cloud Formation, Heat, etc. allow the details of an infrastructure to be codified into a configuration file. The configuration files allow details of an infrastructure to be codified into a configuration file. The configuration files allow the infrastructure files to be elastically created, modified, and destroyed. Terraform has been inspired by the problems they solve.
2. Configuration Enhancements Offered by Terraform: Terraform similarly uses configuration files to detail the infrastructure setup, but it goes further by being both cloud-agnostic and enabling multiple providers and services to be combined and composed.
3. Orchestration across Providers and Services: For example, Terraform can be used to orchestrate an AWS and an OpenStack cluster simultaneously, while enabling 3rd party providers like Cloudflare and DNSimple to be integrated to provide CDN and DNS services.
4. Unified Management of Entire Infrastructure: This enables Terraform to represent and manage the entire infrastructure with its supporting services, instead of only the subset that exists within a single provider. It provides a single unified syntax, instead of requiring operators to use independent and non-interoperable tools for each platform and service.
5. Separation of Planning from Execution: Terraform also separates the planning phase from the execution phase, by using the concept of an



execution plan. By running *terraform plan*, the current is refreshed and the configuration is consulted to generate an action plan.

6. Terraform Plan and Graph: The plan includes all actions to be taken; which resources will be created, destroyed, and modified. It can be inspected by operators to ensure that it is exactly what is expected. Using *terraform graph*, the plan can be visualized to show dependent ordering. Once the plan is captured, the execution phase can be limited to only the actions in the plan.
7. Lack of Separation between Planning and Execution: Other tools combine the planning and the execution phases, meaning that operators are forced to mentally reason about the effects of a change, which quickly becomes intractable in large infrastructures. Terraform lets operators apply changes with confidence, as they will know exactly what will happen beforehand.

Terraform vs. Boto, Fog, etc.

1. Direct Cloud Provider Access API: Libraries like Boto, Fog, etc. are used to provide native access to cloud providers and services by using their APIs. Some libraries are focused on specific clouds, while others attempt to bridge them all and mask the semantic differences.
2. Higher Level Infrastructure Support Tooling: Using a client-library only provides low-level access to APIs, requiring application developers to create their own tooling to build and manage their infrastructure.
3. Syntactical Management of Cloud Infrastructure: Terraform is not intended to give low-level programmatic access to providers, but instead provides a high-level of syntax for describing how cloud resources and services should be created, provisioned, and combined.



4. Uniformness across Providers and Provisioners: Terraform is very flexible, using a pug-in based model to support providers and provisioners, giving it the ability to support almost any service that exposes the APIs.

Terraform vs. Custom Solutions

1. Home-grown Tooling Infrastructure Mechanism: Most organizations start by manually managing infrastructure through simple scripts or web-based interfaces. As then infrastructure grows, any manual approach to management becomes both error-prone and tedious, and many organizations begin to home-rill tooling to help automate the mechanical processes involved.
2. Building and Maintaining Tool Suite: The tools require time and resources to build and maintain. As tools of necessity, they represent the minimum viable features needed by an organization, being built to handle only the immediate needs. As a result, they are often hard to extend and difficult to maintain. Because the tooling must be updated in lockstep with any new features or infrastructure, it becomes the limiting factor for how quickly the infrastructure can evolve.
3. Unified, Simple Resource Management Syntax: Terraform is designed to tackle these challenges. It provides a simple, unified syntax allowing almost any resource to be managed without learning new tooling. By capturing all the resources required, the dependencies between them can be resolved automatically so that operators do not need to remember and reason about them. Removing the burden of building the tool allows the operators to focus on their infrastructure and not the tooling.
4. Open-Source Nature of Terraform: Furthermore, terraform is an open source tool. In addition to HashiCorp, the community around Terraform helps extend its features, fix bugs, and document new use cases. Terraform



helps solve a problem that exists in every organization and provides a standard that can be adopted to avoid reinventing the wheel between and within organizations. The open-source nature ensures that it will be around in the long term.



Sample Terraform Configurations

Overview

1. Using the Samples in this Chapter: The examples in this chapter illustrate some of the ways Terraform can be used. All samples are ready to run as-is. Terraform will ask for input of things such as variables and API keys. To continue using the sample, the parameters should be saved in **terraform.tvvars** file or in a *provider* configuration block.
2. Caveat behind the Sample Usage: The samples real providers that launch *real* resources. That means they can cost money to experiment with. To avoid any unexpected changes, one must be sure to understand the price of the resources before launching them, and verify that any unneeded resources are cleaned up afterwards.

Samples

1. Repos Containing all the Samples: The samples are distributed across several repos. The README file in the Terraform repo has links to all of them (<https://github.com/hashicorp/terraform/tree/master/examples>).
2. Local Installation of Terraform: To use these samples, Terraform must first be installed on the local machine. It may be installed from the downloads page (/downloads.html). Once installed, the samples may be downloaded, viewed, and run.



3. Cloning the Sample from Github: To use a sample, the repo that contains it must be cloned and the corresponding directory navigated to. For example, to try the AWS two-tier architecture sample:

```
git clone https://github.com/terraform-providers/terraform-provider-aws.git
cd terraform-provider-aws/examples/two-tier
```

4. Sequential Execution of Terraform Commands: Any preferred code editor may be used to browse and read the configurations. To try out an example, Terraform's *init* and *apply* commands must be run while in the example's directory:

```
terraform init
:
terraform apply
```

Terraform will interactively ask for the variable input and potentially the provider configuration, and will start executing. Once the example is done with, *terraform destroy* must be run to clean up.

Two-Tier AWS Architecture

1. Github Location: <https://github.com/terraform-providers/terraform-provider-aws/tree/master/examples/two-tier>
2. Two-tier AWS Architecture Template: This provides a template for running a simple two-tier architecture for Amazon Web Services. The premise is that you have stateless app servers running behind an ELB serving traffic.



3. Deploying Application to the Servers: To simplify the example, this intentionally ignores deploying and getting the application onto the servers. However, this could be done so either via provisioners (</docs/provisioners/index.html>) and a configuration management tool, or by pre-baking configured AMIs with Packer (<https://www.packer.io>).
4. AWS Region EC2 Key Pair Creation: The example will also create a new EC2 Key Pair inside the specified AWS Region. The key name and the path to the public key must be specified via the Terraform command *vars*.
5. Provisioning Component inside the Script: After *terraform apply* is run on this configuration, it will automatically output the DNS address of the ELB. After the instance registers, this should respond with the default *nginx* web page.
6. Configuring AWS and Running Terraform: To run, the AWS provider is configured as described in <https://www.terraform.io/docs/providers/ws/index.html>. A command such as the following is used to run:

```
terraform apply -var 'key_name  
= {aws_key_name}' -var 'public_key_path  
= {location_of_key_in_local_machine}'
```

For example:

```
terraform apply -var 'key_name = terraform' -var 'public_key_path  
= /Users/jsmith/.ssh/terraform.pub'
```

Alternative to using *-var* with each command, the **terraform.template.tfvars** file can be copied to **terraform.tfvars** and updated.



Terraform-provider-aws/examples/two-tier/terraform-template.tf

```
key_name = "terraform-provider-aws-example"
public_key_path = "~/.ssh/terraform-provider-aws-example.pub"
```

Terraform-provider-aws/examples/two-tier/outputs.tf

```
output "address" {
  value = aws_elb.web.dns_name
}
```

Terraform-provider-aws/examples/two-tier/main.tf

1. Minimum Acceptable Terraform Version:

```
terraform {
  required_version = " ≥ 0.12"
}
```

2. Setting the Provider Region:

```
provider "aws" {
  region = var.aws_region
}
```

3. VPC to Launch the Instances into:



```
resource "aws_vpc" "default" {  
    cidr_block = "10.0.0.0/16"  
}
```

4. Internet Gateway to give Subnet access to the Outside World:

```
resource "aws_internet_gateway" "default" {  
    vpc_id = aws_vpc.default.id  
}
```

5. VPC Internet Access on the Main Route Table:

```
resource "aws_route" "internet_access" {  
    route_table_id = aws_vpc.default.main_route_table_id  
    destination_cidr_block = "10.0.0.0/0"  
    gateway_id = aws_internet_gateway.default.id  
}
```

6. Subnet to Launch the Instances into:

```
resource "aws_subnet" "default" {  
    vpc_id = aws_vpc.default.id  
    cidr_block = "10.0.1.0/24"  
    map_public_ip_on_launch = true  
}
```

7. Security Group for the ELB so that it is accessible through the Web:

```
resource "aws_security_group" "elb" {
```



```
name = "terraform_example_elb"
description = "Used in the Terraform"
vpc_id = aws_vpc.default.id

# HTTP Access from the Web
ingress {
    from_port = 80
    to_port = 80
    protocol = "tcp"
    cidr_blocks = ["10.0.0.0/16"]
}

# Outbound Internet Access
egress {
    from_port = 0
    to_port = 0
    protocol = "-1"
    cidr_blocks = ["0.0.0.0/0"]
}
}
```

8. Setup of the Provider ELB:

```
resource "aws_elb" "elb" {
    name = "terraform_example_elb"
    subnets = [aws_subnet.default.id]
    security_groups = [aws_security_group.elb.id]
    instances = [aws_instance.web.id]

# Setting the Provider ELB Listener
```



```
listener {  
    instance_port = 80  
    instance_protocol = "http"  
    lb_port = 80  
    lb_protocol = "http"  
}  
}
```

9. Setting up the AWS Key Pair Authorization:

```
resource "aws_key_pair" "auth" {  
    key_name = var.key_name  
    public_key = file(var.public_key_path)  
}
```

10. Setting up the AWS Web Instance:

```
resource "aws_instance" "web" {  
    # Setting the AWS Instance Type  
    instance_type = "t2.micro"  
  
    # AMI Based on the Region specified  
    ami = var.aws_amis[var.aws_region]  
  
    # Name of the SSH Key-pair created above  
    key_name = aws_key_pair.auth.id  
  
    # Security Group to allow HTTP and SSH Access  
    vpc_security_group_ids = [aws_security_group.default.id]
```



```
# Launch occurs in the same instance as ELB. In a production
environment it is more common to have a separate private subnet for
backend instances
subnet_id = aws_subnet_default.id

# The connection block tells the provisioner how to communicate with
the resource instance
connection {
    type = "ssh"

    # The default username for the AMI
    user = "ubuntu"
    host = self.public_ip

    # The connection will use the local SSH agent for authentication
}

# The remote provisioner is run on the instance after creating it. This
case just installs nginx and starts it. By default, this should be on port
80.
provisioner "remote-exec" {
    inline [
        "sudo apt - get -y update",
        "sudo apt - get -y install nginx",
        "sudo service nginx start",
    ]
}
}
```

Terraform-provider-aws/examples/two-tier/variables.tf



1. Setting up the AWS Web Instance:

```
variable "public_key_path" {  
    description = << DESCRIPTION Path to the SSH public key to be used for  
    authentication. This key-pair must be added to the local SSH agent so  
    provisioners can connect. Example: ~/.ssh/terraform.pub DESCRIPTION  
}
```

2. Specification of the Public Key Name:

```
variable "key_name" {  
    description = "Desired Name of the AWS Key Pair"  
}
```

3. Specification of the AWS Region:

```
variable "aws_region" {  
    description = "AWS region to launch server"  
    default = "us - west - 2"  
}
```

4. Specification the Ubuntu AWS AMIs: Ubuntu Precise 12.04 LTS (x64)

```
variable "aws_amis" {  
    default = {  
        eu - west - 1 = "ami - 674cbc1e"  
        us - east - 1 = "ami - 1d4e7a66"  
        us - west - 1 = "ami - 969ab1f6"  
        us - west - 2 = "ami - 8803e0f0"
```



```
}  
}
```

Cross Provider Example

1. Github Location:

<https://github.com/hashicorp/terraform/tree/master/examples/cross-provider>

2. What is Cross Provider? This is a simple example of the cross-provider capabilities of Terraform. It creates a Heroku application and points a DNS CNAME record at the result via DNSimple. A *host* query to the outputted hostname should reveal the correct DNS configuration.

Terraform-provider-aws/examples/count/variables.tf

```
variable "aws_region" {  
    description = "AWS region to launch server"  
    default = "us - west - 2"  
}
```

Terraform-provider-aws/examples/count/outputs.tf

```
output "address" {  
    value = aws_elb.web.dns_name  
}
```




Terraform-provider-aws/examples/count/main.tf

1. Required Terraform Version:

```
terraform {  
    required_version = " ≥ 0.12"  
}
```

2. Setting the Provider Region:

```
provider "aws" {  
    region = var.aws_region  
}
```

3. Web Instance of AWS ELB:

```
resource "aws_elb" "elb" {  
    name = "terraform - example - elb"  
  
    # Same availability zone as the instances  
    availability_zones = aws_instance.web.*.availability_zone  
  
    # The instances are registered automatically  
    instances = aws_instance.web.*.id  
  
    # AWS ELB Listener Settings  
    listener {  
        instance_port = 80  
        instance_protocol = "http"  
        lb_port = 80
```



```
        lb_protocol = "http"
    }
}
```

4. AWS AMI Settings:

```
data "aws_ami" "ubuntu" {
    most_recent = true
    owners = [099720109477] # Canonical

    # The instances are registered automatically
    instances = aws_instance.web.*.id

    # AMI Filter Settings #1
    filter {
        name = "name"
        values = ["ubuntu/images/hvm-ssd/ubuntu-focal-20.04-
                  amd64-server-*"]
    }

    # AMI Filter Settings #1
    filter {
        name = "virtualization type"
        values = ["hvm"]
    }
}
```

5. Setting the Provider Region:

```
resource aws_instance "web" {
```



```
instance_type = "t2.small"
ami = data.aws_ami.ubuntu.id

# This will create 4 instances
count = 4
}
```

Consul Example

1. Github Location:
<https://github.com/hashicorp/terraform/tree/master/examples/consul-provider>
2. What is the Consul Tool? Consul (<https://www.consul.io>) is a tool for service discovery, configuration, and orchestration. The key/value store it provides is often used to store the application configuration and information about the infrastructure necessary to process requests.
3. Interfacing with the Consul Provider: Terraform provides a Consul provider ([/docs/providers/consul/index.html](https://www.terraform.io/docs/providers/consul/index.html)) which can be used to interface with Consul from inside a Terraform configuration.
4. EC2 Configuration in Demo Consul: This example uses the Consul demo cluster (<https://demo.consul.io>) to both read configuration and store information about a newly created EC2 instance. The size of the EC2 instance will be determined by the `tf_test/size` key in Consul, and will default to `m1.small` if that key does not exist. Once the instance is created, `tf_test/id` and `tf_test/public_dns` keys will be set to the computed values for the instance.
5. Initial Setting of AWS Configuration: Before running the example, the Web UI <https://demo.consul.io/ui/dc1/kv/> is used to set the `tf_test/size` key to `t1_micro`. Once that is done, the configuration is copied into a configuration



file - *consul.tf* works fine. Either the AWS credentials are provided as a default value in the configuration or *apply* is invoked with the appropriate variables set.

6. Viewing the Configuration Key Settings: Once *apply* has completed, the keys in the Consul can be seen by visiting the Web UI <https://demo.consul.io/dc1/kv>. The *tf_test/id* and the *tf_test/public_dns* can be seen to have been set.
7. Delete Property for Key Cleanup: The infrastructure can then be torn down (https://learn.hashicorp.com/tutorials/aws-destroy?in-terraform/aws-get-started&utm_source=WEBSITE&utm_medium=WEB_IO&utm_offer=ARTICLE_PAGE&utm_content=DOCS). Because the *delete* property of two of the Consul keys are set, Terraform will cleanup up those keys on destroy.
8. Decoupling Infrastructure Details from Architecture: Inputs like AML names, security groups, puppet roles, bootstrap scripts etc. can all be loaded from Consul. This allows specifics of an infrastructure to be decoupled from its overall architecture. This enables details to be changed without updating the Terraform configuration.
9. Storing Resource Attributes in Consul: Outputs from Terraform can also be easily stored in Consul. One powerful feature this enables is using Consul for inventory management. If an application relies on ELB for routing, Terraform can update the application's configuration directly by setting the ELB address directly into Consul. Any resource attribute can be stored in Consul, allowing an operator to capture anything useful.

Terraform-provider-consul/examples/kv/variables.tf

1. AWS Region to Create Resources in:

```
variable "aws_region" {
```



```
description = "The AWS region to create resources in"
default = "us - east - 1"
}
```

2. AWS AMI Image Providers:

```
# AMI's from http://cloud-images.ubuntu.com/locator/ec2
variable "aws_amis" {
  default = {
    eu - west - 1 = "ami - 674cbc1e"
    us - east - 1 = "ami - 1d4e7a66"
    us - west - 1 = "ami - 969ab1f6"
    us - west - 2 = "ami - 8803e0f0"
  }
}
```

Terraform-provider-consul/examples/kv/main.tf

1. Consul Provisioner for the Demo Cluster:

```
provider "consul" {
  address = "demo.consul.io:80"
  datacenter = "nyc1"
}
```

2. Setup an AWS Provider:

```
provider "aws" {
  region = "${var.aws_region}"
}
```



```
}
```

3. Setting Up Consul Keys to Provide Inputs:

```
data "consul_keys" "input" {  
  key {  
    name = "size"  
    path = "tf_test/size"  
    default = "m1.small"  
  }  
}
```

4. New AWS Instance using Dynamic AMI and Instance Type:

```
resource "aws_instance" "test" {  
  ami = "${lookup(var.aws_amis, var.aws_region)}"  
  instance_type = "${data.consul_keys.input.var.size}"  
}
```

5. Consul Key stores Instance ID and DNS Name of the Instance:

```
resource "consul_keys" "test" {  
  key {  
    name = "id"  
    path = "tf_test/id"  
    value = "${aws_instance.test.id}"  
    delete = true  
  }  
  key {  
    name = "address"
```



```
path = "tf_test/public_dns"  
value = "${aws_instance.test.public_dns}"  
delete = true  
}  
}
```



Kubernetes

Overview

1. Focus of the Kubernetes System: *Kubernetes* is an open-source container orchestration system for automating computer application deployment, scaling, and management (Garrison (2016), Wikipedia (2020)).
2. Objectives of the System Design: It aims to provide a platform for automating deployment, scaling, and operations of application containers across clusters of hosts. It works with a range of container tools, including Docker.
3. Vendor Service Deployments using Kubernetes: Many closed services offer a Kubernetes-based platform or infrastructure as a service (PaaS or IaaS) on which Kubernetes can be deployed as a platform providing service. Many vendors also provide their own branded Kubernetes distributions.

Kubernetes Objects

1. Purpose behind the Kubernetes Objects: Kubernetes maintains a set of building blocks or primitives which collectively provide mechanisms that deploy, maintain, and scale applications based on CPU, memory, or custom metrics (Sharma (2017)).
2. Structure and Extensibility of Kubernetes: Kubernetes is loosely coupled and extensible to meet different workloads. The extensibility is provided in large part by the Kubernetes API, which is used by internal components as well as extensions and containers that run on Kubernetes. The platform



exerts its control over compute and storage resources by defining resources as Objects, which can then be managed as such.

Pods

1. What is a Kubernetes Pod? A *pod* is a higher level of abstraction grouping containerized components. A pod consists of one or more containers that are guaranteed to be co-located on the host machine and can share resources. The basic scheduling unit in Kubernetes is a *pod*.
2. IP Addresses for Kubernetes Pods: Each pod in Kubernetes is assigned a unique *Pod IP address* within the cluster, which allows applications to use ports without the risk of conflict (Langemak (2015a)). Within the pod, all containers can reference each other on localhost, but a container within one pod has no way of directly addressing another container within another pod; for that, it has to use the pod IP address.
3. Name Service to Reference Pods: An application developer should never use the pod IP address though, to reference/invoke a capability in another pod, as pod IP addresses are ephemeral - the specific pod that they are addressing may be assigned to another pod IP address on re-start. Instead, they should use a reference to a Service, which holds a reference to the target pod at the specific pod IP address.
4. Local/Network Disk Volume Pods: A pod can define a Volume, such as a local disk directory or a network disk, and expose it to the containers in the pod (Strachan (2015a)). Pods can be managed manually through the Kubernetes API, or their management can be delegated to a controller.
5. Storage for ConfigMaps and Secrets: Such volumes are also the basis for Kubernetes features of ConfigMaps - which provide access to configuration through the file-system visible to the container - and Secrets - which provide access to credentials required to access remote resources securely



by providing those credentials on the filesystem visible only to authorized containers.

ReplicaSets

1. Purpose of the ReplicaSet: A ReplicaSet's purpose is to maintain a stable set of replica pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical pods.
2. Selector Based ReplicaSet Specification: The ReplicaSets can also be said to be a grouping mechanism that lets Kubernetes maintain the number of instances that have been declared for a given pod. The definition of ReplicaSet uses a selector whose evaluation will result in identifying all pods that are associated with it.

Services

1. Definition of a Kubernetes Service: A Kubernetes source is a set of pods that work together, such as one-tier of a multi-tier application. The set of pods that constitute a service is defined by a label selector.
2. Modes of Service Discovery in Kubernetes: Kubernetes provides two modes of service discovery, using environmental variables of Kubernetes DNS. Service discovery assigns a stable IP address and DNS name to the service, and load balances traffic in a round-robin manner to network connections of that IP address among the pods matching the selector - even as failures cause the pods to move from machine to machine.
3. Exposing Front/Back End Services: By default, a service is exposed inside a cluster, e.g., back end pods might be grouped into a service, with requests from the front-end pods load balanced among them, but a service can also



be exposed outside a cluster, e.g., for clients to reach front-end pods (Langemak (2015b)).

Volumes

1. Transient Nature of Kubernetes Filesystems: Filesystems in Kubernetes containers provide ephemeral storage, by default. This means that re-start of the pod will wipe out any data on such containers, and therefore, this form of storage is quite limiting in anything but trivial applications.
2. Persistent Storage Mechanism - Kubernetes Volume: A Kubernetes Volume provides persistent storage that exists for the life of the pod itself. This storage can also be used as shared disk space for containers within the pod.
3. Restrictions on the Mount Points: Volumes are mounted at specific mount points within the container, which are defined by the pod configuration, and cannot mount onto other Volumes or link to other Volumes. The same Volume can be mounted at different points in the filesystem tree by different containers.

Namespaces

Kubernetes provides a partitioning of the resources it manages into non-overlapping sets called namespaces. They are intended for use in environments with many users spread across multiple teams or projects, even separating environments like development, test, and production.

ConfigMaps and Secrets



1. Types of Kubernetes Configuration Data: A common application challenge is deciding where to store and manage configuration information, some of which may contain sensitive data. Configuration data can be anything as fine-grained as individual properties or coarse-grained information like entire configuration files or JSON/XML documents.
2. Access to Configuration Data: Kubernetes provides two closely related mechanisms to deal with this need - *configmaps* and *secrets*, both of which allow for configuration changes to be made without requiring an application rebuild. The data from configmaps and secrets will be made available to every single application to which these objects have been bound via deployment.

StatefulSets

1. Kubernetes Stateless versus Stateful Scaling: It is very easy to address the scaling of stateless applications: one simply adds more running pods - which is something that Kubernetes does very well. Stateful workloads are much harder, because the state needs to be preserved if a pod is restarted, and if the application is scaled up or down the state may need to be redistributed.
2. Ordering of the Stateful Instances: Databases are an example of stateful workloads. When run in high-availability mode, many databases come with the notion of a primary instance and secondary instance(s). In this case, the notion of ordering the instances is important.
3. Uniqueness of the Stateful Instances: Other applications like Kafka distribute the data among their brokers - so that one broker is not the same as another. In this case, the notion of instance uniqueness is important.
4. Uniqueness and Ordering among Pod Instances: StatefulSets are controllers that are provided by Kubernetes that enforce the properties of uniqueness



and ordering among the instances of a pod and can be used to run stateful applications.

DaemonSets

Normally, the location where the pods are run are determined by the algorithm implemented in the Kubernetes scheduler. The ability to do this kind of pod scheduling is implemented by the feature called DaemonSets.

Secrets

Secrets contain ssh keys, passwords, and OAuth tokens for the pod.

Managing Kubernetes Objects: Labels and Selectors

1. Label Selector Based Object Matches: Kubernetes enables clients - users or internal components - to attach keys called *labels* to any API Object in the system, such as pods or nodes. Correspondingly, *label selectors* are queries against labels that resolve to matching objects.
2. Service Routing using Label Selectors: When a service is defined, one can define the label selector that will be used by the service router/load balancer to select the pod instances that the traffic will be routed to.
3. Controlling Service Routing Using Labels: Thus, simply changing the labels of the pods or changing the label selectors on the service can be used to control which pod gets the traffic and which doesn't, and this can be used to support various deployment patterns like blue-green deployments or A-B



testing. This capability to dynamically control how services utilize implementing resources provides a loose coupling within the infrastructure.

4. Multi-Criterion Label Selection Example: For example, if the application's pods have labels for a system *tier* - with values such as *front_end*, *back_end*, for example - and a *release_track* - with values such as *canary*, *production*, for example - then an operation on all of *back_end* and *canary* nodes can use a label selector, such as:

$$tier = back_end \text{ AND } release_track = canary$$

Field Selectors

1. Object-Attribute-Value Based Selection: Just like fields, field selectors also let one select Kubernetes resources. Unlike labels, the selection is based on the attribute values inherent to the resource being selected, rather than user-defined categorization.
2. Built-in Metadata Field Selectors: *metadata.name* and *metadata.namespace* are field selectors that will be present on all Kubernetes objects. Other selectors can be used depending on the object/resource type.

Replication Controllers and Deployments

1. Focus of the Replication Controller: A ReplicaSet declares the number of instances of a pod that is needed, and the Replication Controller manages the system so that the number of healthy pods that are running matches the number of pods declared in the ReplicaSet, determined by evaluating its selector.



2. Kubernetes Deployments versus ReplicaSets: Deployments are a higher-level management mechanism for ReplicaSets. While the Replication Controller manages the scale of the ReplicaSet, Deployments will manage what happens to the ReplicaSet - whether an update has to be rolled out, rolled back, etc.
3. Deployment Induced Changes in ReplicaSets: When Deployments are scaled up or down, this results in the declaration of the ReplicaSet changing - and this change in the declared state is managed by the Replication Controller.

Cluster API

1. Programmatic Control of Kubernetes Clusters: The design principles underlying Kubernetes allow one to programmatically create, configure, and manage Kubernetes clusters. This function is exposed via an API called the Cluster API.
2. Cluster as a Controllable Resource: A key concept embodied in the API is the notion that the Kubernetes cluster itself is a resource object that can be managed just like any other Kubernetes resources. Similarly, machines that make up the cluster are also treated as a Kubernetes resource.
3. Cloud Provider Specific API Functions: The API has two pieces - the core API, and a provider implementation. The provider implementation consists of cloud-provider specific functions that let Kubernetes provide the cluster API in a fashion that is well-integrated with the cloud provider's services and resources.

Architecture



Kubernetes follows a primary/replica architecture. The components of Kubernetes can be divided into those that manage an individual node, and those that are part of the control plane.

Kubernetes Control Plane

1. Kubernetes Master and Control Plane: The Kubernetes master is the main controlling unit of the cluster, managing the workload and directing communication across the system. The Kubernetes control plane consists of various components, each its own process, that can run both on a single master node or on multiple masters supporting high-availability clusters. The various components of the Kubernetes control plane are as follows.
2. etcd - Distributed Key-Value Store: etcd is a persistent, lightweight, distributed key-value store developed by CoreOS that reliably stores the configuration data of the cluster, representing the overall state of the cluster at any given point in time.
3. Choosing Cluster Consistency over Availability: Just like Apache ZooKeeper, etcd is a system that favors consistency over availability in the event of network partition. This consistency is crucial for correctly scheduling and operating services.
4. Usage of etcd's Watch API: The Kubernetes API Server uses etcd's Watch API to monitor the cluster or roll out critical configuration changes or simply restore any divergences of the state of the cluster back to what was declared by the deployer.
5. Example - Creating Additional Pod Instance: As an example, if the deployer specified that three instances of a particular pod need to be running, this fact is stored in etcd. If it is found that only two instances are running, this delta will be detected by comparison with the etcd data, and Kubernetes will use this to schedule the creation of an additional instance of that pod.



6. The Kubernetes API Server Component: The API Server is a key component and serves the Kubernetes API using JSON over HTTP, which provides both the internal and the external interfaces to Kubernetes (Marhubi (2015a)). The API server process and validates REST requests and updates the state of the API objects in etcd, thereby allowing clients to configure workloads and containers across worker nodes (Ellingwood (2018)).
7. Role of the Kubernetes Scheduler: The scheduler is the pluggable component that selects which node an unscheduled pod - the basic entity managed by the scheduler - runs on, based on resource availability. The scheduler tracks resource use on each node to ensure that the workload is not scheduled in excess of available resources.
8. Matching Resource Supply and Workload Demand: For this purpose, the scheduler must know the resource requirements, resource availability, and other user-provided constraints and policy directives such as quality-of-service, affinity/anti-affinity requirements, data locality, and so on. In essence, the scheduler's role is to match the resource *supply* to workload *demand*.
9. The Kubernetes Controller Manager Component: The controller is a reconciliation loop that drives the actual cluster state toward the desired cluster state, communicating with the API Server to create, update, and delete the resources it manages, i.e., the pods, the services, and the end-points (Marhubi (2015a)). The Controller Manager is a component that manages a set of core Kubernetes controllers.
10. The Kubernetes Replication Controller Component: One kind of controller is the Replication Controller, which handles replication and scaling by running a specified number of pods across the cluster. It also handles creating replacement pods if the underlying node fails.
11. Kubernetes DaemonSet and Job Controllers: Other controllers that are a part of the core Kubernetes system include a DaemonSet Controller for



running exactly one pod on every machine, and a Job Controller for running pods that run to completion, e.g., as part of a batch job (Sanders (2015)).

12. Control Exercised using Label Selection: The set of pods that a controller manages is determined by the label selection that are part of the controller's definition.

Kubernetes Node

1. Nodes where Containers are Deployed: A Node, also known as a Worker or a Minion, is a machine where containers/workloads are deployed. Every node in the cluster must run a container runtime such as Docker, as well as the below-mentioned components, for communication with the primary for network configuration of these containers.
2. Kubelet - Kubernetes Node Health Monitor: Kubelet is responsible for running the state of each node, ensuring that all containers in the node are healthy. It takes care of starting, stopping, and maintaining application containers organized into pods as directed by the control plane (Marhubi (2015b)).
3. Mechanism Employed in Monitoring the Node: Kubelet monitors the state of a pod, and if not in the desired state, the pod re-deploys to the same node. Node status is relayed every few seconds via heartbeat messages to the primary. Once the primary detects a node failure, the Replication Controller observes this state change and launches pods on other healthy nodes.
4. Role of the Kube-Proxy: The Kube-Proxy is an implementation of the network proxy and a load balancer, and it supports service abstraction along with other networking operations. It is responsible for routing the traffic to the appropriate container based on the IP and the port number of the incoming request.



5. Kubernetes Workload/Container Runtime: A container resides inside a pod. The container is the lowest level of a micro-service, which holds the running applications, libraries, and their dependencies. Containers can be exposed to the world through an external IP address. Kubernetes supports Docker containers since its first version, and in July 2016 rkt container engine was added.

Add-Ons

1. Additional Special Purpose Cluster Applications: Add-ons operate just like any other application running within the cluster: they are implemented via pods and services, and are only different in that they implement features of the Kubernetes cluster. The pods may be managed by Deployments, Replication Controllers, and so on. There are many add-ons, and the list is growing. Some of the more important are as follows.
2. DNS: All Kubernetes clusters should have a cluster DNS; it is a mandatory feature. Cluster DNS is a DNS server, in addition to the other DNS server(s) in the environment, which serves DNS records for Kubernetes services. Containers started by Kubernetes automatically include this DNS server in their DNS searches.
3. Web UI: This is a general purpose, web-based UI for Kubernetes clusters. It allows users to manage and troubleshoot applications running in the cluster, as well as the cluster itself.
4. Container Resource Monitoring: Providing a reliable application runtime, and being able to scale it up or down in response to workloads, means being able to continuously and effectively monitor workload performance. Container Resource Monitoring provides this capability by recording metrics about the containers in a central database, and provides UI for browsing that data. The cAdvisor is a component of the slave node that provides a



limited metric monitoring capability. There are full metrics pipelines as well, such as Prometheus, which can meet most monitoring needs.

5. Cluster-Level Logging: Logs should have a separate storage and lifecycle independent of nodes, pods, or containers. Otherwise, node or pod failures can cause loss of event data. The ability to do this is called cluster-level logging, and such mechanisms are responsible for saving container logs to a central log-store with searching/browsing interface. Kubernetes provides no native storage for log data, but one can integrate many existing logging solutions into the Kubernetes cluster.

Microservices

Kubernetes is commonly used as a way to host microservice-based implementation, because it and its associated ecosystem of tools provides all the capabilities needed to address key concerns of any microservice architecture.

Kubernetes Persistence Architecture

1. Containers - Enabling Portability of Software: Containers emerged as a way to make software portable. The container contains all the packages you need to run a service. The provided filesystem makes the containers extremely portable and easy to use in development. A container can also be moved from development to test or production with no or relatively few changes.
2. Reasons for Requiring Container Persistence: Historically, Kubernetes was suitable only for stateless services. However, many applications have a database, which require persistence, which leads to creation of persistence



storage for Kubernetes. Implementing persistent storages is one of the top challenges for Kubernetes administrators, DevOps, and cloud engineers. Containers may be more ephemeral, but more and more of their data is not, so one needs to ensure data's survival in case of container termination or hardware failure.

3. Persistent Storage for Containerized Applications: When deploying containers with Kubernetes or containerized applications, companies often realize that they need persistent storage. They need to provide fast and reliable storage for databases, root images, and other data used by the containers.
4. CNCF Kubernetes Persistent Storage Guidelines: In addition to the landscape, the Cloud Native Computing Foundation (CNCF) has published other information about Kubernetes Persistent Storage including a blog helping to define the container attached storage pattern. This pattern can be thought of as one that uses Kubernetes itself as a component of the storage system or service.
5. Public Domain Storage Orchestration Projects: More information about the relative popularity of these and the other approaches can be found on the CNCF's landscape survey as well, which showed that OpenEBS from MayaData and Rook - a storage orchestration project - were the two projects most likely to be in evaluation.

References

- Ellingwood, J. (2018): [An Introduction to Kubernetes](#)
- Garrison, J. (2016): [Why Kubernetes is Abbreviated k8s](#)
- Langemak, J. (2015a): [Kubernetes 101 - Networking](#)
- Langemak, J. (2015b): [Kubernetes 101 - External Access into the Cluster](#)
- Marhubi, K. (2015a): [Kubernetes from the Ground Up: the API Server](#)



- Marhubi, K. (2015b): [What even is a kubelet?](#)
- Sanders, J. (2015): [Kubernetes: Exciting Experimental Features](#)
- Sharma, P. (2017): [Autoscaling based on CPU/Memory in Kubernetes — Part II](#)
- Strachan, J. (2015): [Kubernetes for Developers](#)
- Wikipedia (2020): [Kubernetes](#)



Apache Kafka

Overview

1. Stream-Processing Low-Latency Platform: *Apache Kafka* is an open-source stream-processing software developed by Apache Software Foundation, written in Java and Scala. The project aims to provide a unified high-throughput, low-latency platform for handling real-time data feeds.
2. Data Interfacing and Streams Library: Kafka can connect to external systems for data import/export via Kafka Connect and provides Kafka Streams, a Java Streams processing library (Wikipedia (2020)).
3. TCP Based Bulk Message Abstraction: Kafka uses a binary TCP protocol that is optimized for efficiency and relies on a *message set* abstraction that naturally groups messages together to reduce the overhead of network round trips. This leads to larger network packets, larger sequential disk operations, and contiguous memory blocks which allow Kafka to convert a bursty stream of random message writes into linear writes.

Architecture

1. Kafka Producers, Partitions, and Topics: Kafka stores key-value messages that come from arbitrarily many processors called *producers*. The data is partitioned into different *partitions* within different *topics*.
2. Ordering of Messages within Partitions: Within a partition, messages are strictly ordered by offsets - the position of message within a partition - and



indexed and stored together with a timestamp. Other processors called *consumers* can read messages from partitions.

3. Streams Processing API and Interfaces: For streams processing, Kafka offers the Streams API that allows writing Java applications that consumer data from Kafka and writes results back to Kafka. Apache Kafka also works with external stream processing systems such as Apache Apex, Apache Flink, Apache Spark, Apache Storm, and Apache NiFi.
4. Kafka Brokers and Partition Distribution: Kafka runs on a cluster of one or more servers called brokers, and the partitions of all topics are distributed across the cluster nodes. Additionally, partitions are replicated across multiple brokers.
5. Stream Volumes and Transactional Features: This architecture allows Kafka to deliver massive streams of messages in a fault-tolerant fashion and has allowed it to replace some of the conventional messaging systems like java Messaging Service (JMS), Advanced Message Queue Protocol (AMQP), etc. Since the 0.11.0.0 release, Kafka offers *transactional writes* which provide exactly-once stream-processing using the Streams API.
6. Regular Topics - Purpose and Usage: Kafka supports two types of topics - regular and compacted. Regular topics can be configured with a retention time of space bound. If there are records that are older than the specified retention time or if the space bound is exceeded for a partition, Kafka is allowed to delete old data to free storage space. By default, topics are configured with a retention time of 7 days, but it is also possible to store data indefinitely.
7. Compacted Topics - Purpose and Usage: For compacted topics, records don't expire based on time or space bounds. Instead, Kafka treats the later messages as updates to the older messages with the same key and guarantees never to delete the latest message per key. Users can delete messages entirely by writing a so-called tombstone message with NULL value for a specific key. To this purpose, there are five major APIs in Kafka:



8. Producer API: Permits an application to publish streams of records.
9. Consumer API: Permits an application to subscribe to topics and processes streams of records.
10. Connector API: Executes the reusable producer and consumer APIs that can link the topics to existing applications.
11. Streams API: This API converts the input streams to output and produces the result.
12. Admin API: This API is used to manage Kafka topics, brokers, and other objects.
13. Message Protocol and API/Usage: The consumer and the producer APIs build on top of the Kafka messaging protocols and offer a reference implementation for Kafka consumer and producer clients in Java. The underlying message protocol is a binary protocol that developers can use to write their own consumer or producer clients in any programming language. This unlocks Kafka from the JVM eco-system.

Connect API

1. Handling Data from External Systems: Kafka Connect - or Connect API - is a framework to import/export data from/to other systems. It was added in Kafka 0.9.0.0 release and uses the Producer and the Consumer APIs internally. The Connect Framework itself executes the so-called *connectors* that implement the actual logic to read/write data from other systems.
2. API for Implementing Custom Connectors: The Connect API defines the programming interface that must be implemented to build a custom connector. Many open source and commercial connectors for popular data systems are available already. However, Kafka itself does not include any production ready connection.



Streams API

1. Java Based Streams Processing Library: Kafka Streams - or Streams API - is a stream-processing library written in Java. The library allows for the development of stateful stream-processing applications that are scalable, elastic, and fully fault-tolerant.
2. DSL Standard and Custom Operators: The main API is a stream-processing domain-specific language - DSL - that offers high-level operations like filter, map, grouping, windowing, aggregation, joins and the notion of tables. Additionally, the Processor API can be used to implement custom operators for a more low-level development approach. The DSL and the Processor API can be mixed, too.
3. Stateful Stream Processing using RocksDB: For stateful Streams processing, Kafka Streams uses RocksDB to maintain local operator state. Because RocksDB can write to disk, the maintained state can be larger than the available main memory. For fault-tolerance, all updates to the local state stores are also written into a topic in the Kafka cluster. This allows re-creating the state by reading those topics and feeding all data into RocksDB.

Version Compatibility

Up to version 0.9.x, Kafka brokers are backward compatible with older clients only. Since Kafka 0.10.0.0, brokers are also forward compatible with newer clients. If a newer client connects to an older broker, it can only use the features that the broker supports. For the Streams API, full compatibility starts with version 0.10.0.0; a 0.10.0.0 Kafka Streams application is not compatible with 0.10.0.0 or older nodes.



Performance

Monitoring end-to-end performance requires tracking metrics from brokers, consumers, and producers. In addition to monitoring ZooKeeper, which Kafka uses for coordination among consumers (Mouzakitis (2016)). There are currently several platforms to track Kafka performance, either open-source, like LinkedIn's Burrow (<https://github.com/linkedin/Burrow/wiki>), or paid, like DataDog and Deep BI. In addition to these platforms, collecting Kafka data can also be performed using tools commonly bundled with Java, including JConsole.

References

- Mouzakitis, E. (2016): [Monitoring Kafka Performance Metrics](#)
- Wikipedia (2020): [Apache Kafka](#)



Knowledge Integration

1. Definition and Scope of KI: Knowledge Integration - KI -is the process of synthesizing multiple language models - or representations - into a common model/representation (Wikipedia (2024)).
2. Integration of Knowledge vs. Information: Compared to information integration, which involves merging information having different schemas and representation models, knowledge integration focusses more on synthesizing the understanding of a given subject from different perspectives.
3. Example - KI over Student Grades: For example, multiple interpretations are possible for a set of student grades, typically each from a certain perspective. An overall integrated view and understanding of this information can be achieved if these interpretations can be put under a common model, say, a student performance index.
4. KI as a Knowledge Incorporator: **Knowledge integration** has also been studied as the process of incorporating new information into a body of existing knowledge using an inter-disciplinary approach.
5. Steps involved in the Incorporator Approach: This process involves determining how the new information and the existing knowledge interact, how the existing knowledge should be modified to incorporate the new information, and how the new information should be modified in light of the existing knowledge.
6. Role of Learning Agent: A learning agent that actively investigates the consequences of incorporating the new information can detect and exploit a variety of learning opportunities, e.g., to resolve knowledge conflicts and



the fill knowledge gaps. By incorporating these learning opportunities, the learning agent is able to learn beyond the explicit context of the new information.

7. Expansion of the Knowledge Base: The machine learning program KI, developed at the University of Texas at Austin, was created to study the use of automated and semi-automated knowledge integration to assist knowledge engineers in constructing a large knowledge base.
8. Use of Semantic Matching/Minimal Mapping: A possible technique which can be used in semantic mapping. More recently, a technique used in minimizing mapping validation and visualization has been presented which is based on Minimal Mapping.
9. Approach of the Minimal Mapping Technique: Minimal mappings are high-quality mappings such that:
 - a. All other mappings can be computed from them time linear in the size of the input graphs, and
 - b. None of them can be dropped without losing property a.

References

- Wikipedia (2024): [Knowledge Integration](#)